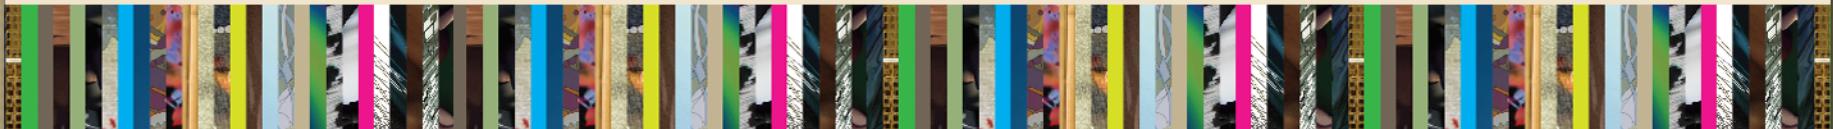




# SIGGRAPH ASIA 2009

革新の波動  
*the pulse of innovation*



# Advanced OpenCL Event Model Usage

Derek Gerstmann

**University of Western Australia**

<http://local.wasp.uwa.edu.au/~derek>



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# OpenCL Event Model Usage

- **Outline**

- Execution Model
- Usage Patterns
- Synchronisation
- Event Model
- Examples



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Execution Model



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Execution Model

- **Command queues are used to submit work to a device**
  - Commands are queued **in-order**
  - Commands execute **in-order** or **out-of-order**



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Execution Model

- **Explicit synchronisation is required for ...**
  - Out-of-order command queues
  - Multiple command queues



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Single Device In-Order Usage Model



- **1x In-Order Queue, 1x Context, 1x Device**
  - Simple and straightforward **in-order** queue
  - All commands execute on single device
  - All memory operations occur in single memory pool



# Single Device In-Order Usage Model

in-order queue

context

CPU

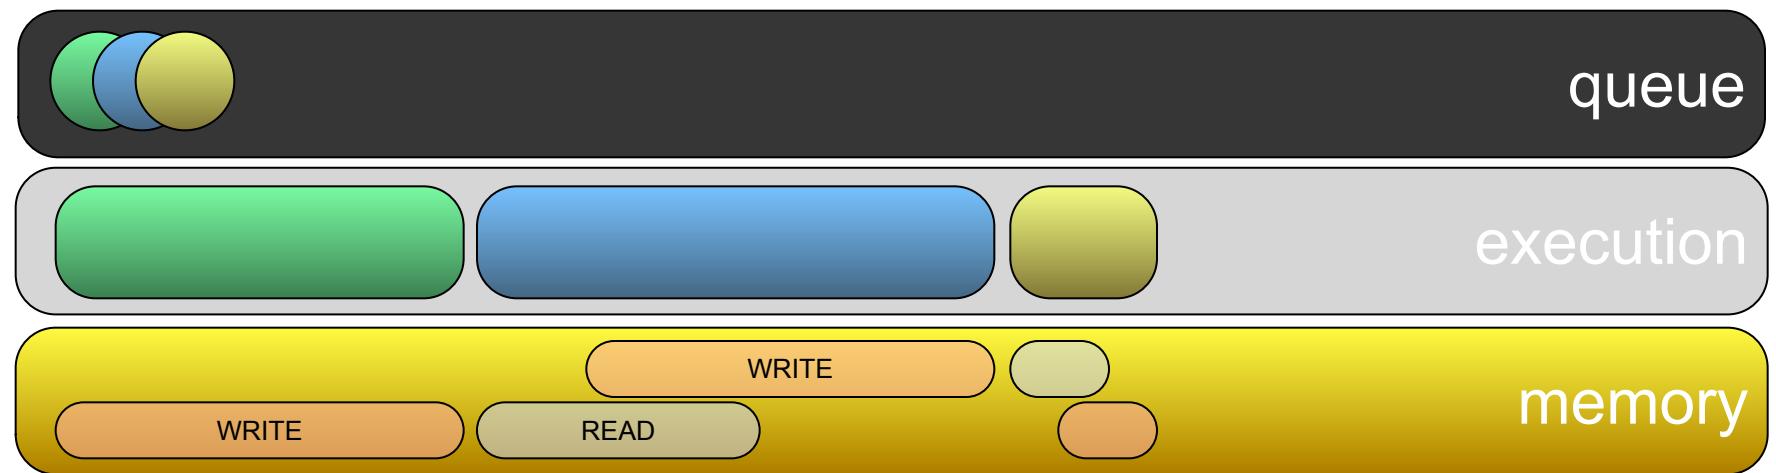
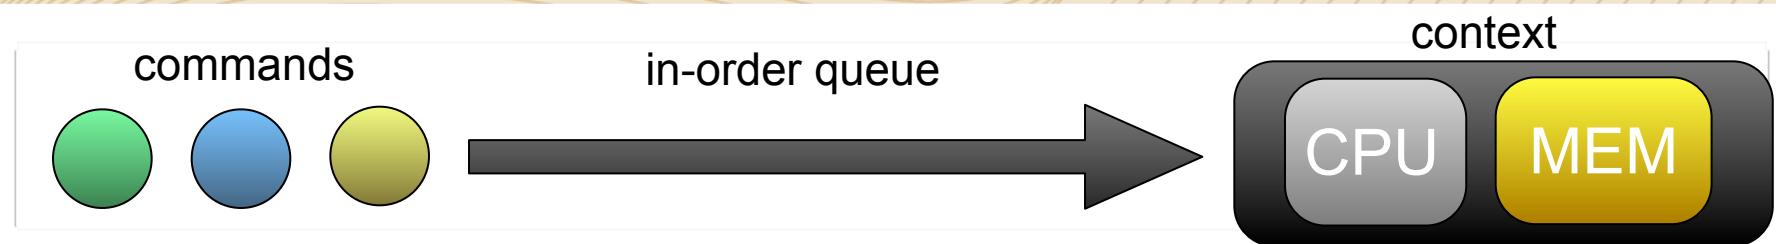
MEM

```
cl_uint num_devices;  
cl_device_id devices[1];  
  
err = clGetDeviceIDs(NULL, CL_DEVICE_TYPE_CPU, 1, &devices[0], &num_devices);  
context = clCreateContext(0, 1, devices, NULL, NULL, &err);  
  
cl_command_queue queue_cpu;  
queue_cpu = clCreateCommandQueue(context, devices[0], 0 /* IN-ORDER */, &err);  
  
/* ... enqueue work ... */
```



SIGGRAPH ASIA 2009  
the pulse of innovation

# Single Device In-Order Usage Model



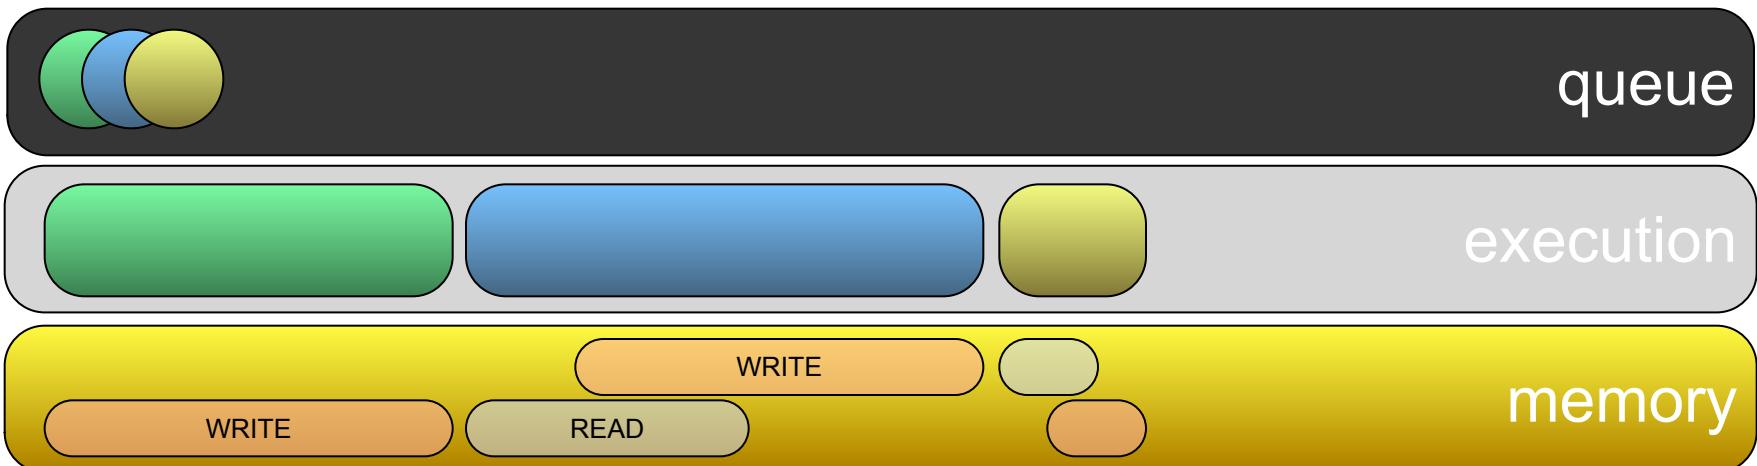
time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Single Device In-Order Usage Model

- Device executes commands after the previous one finishes
- Memory transactions have consistent view



time



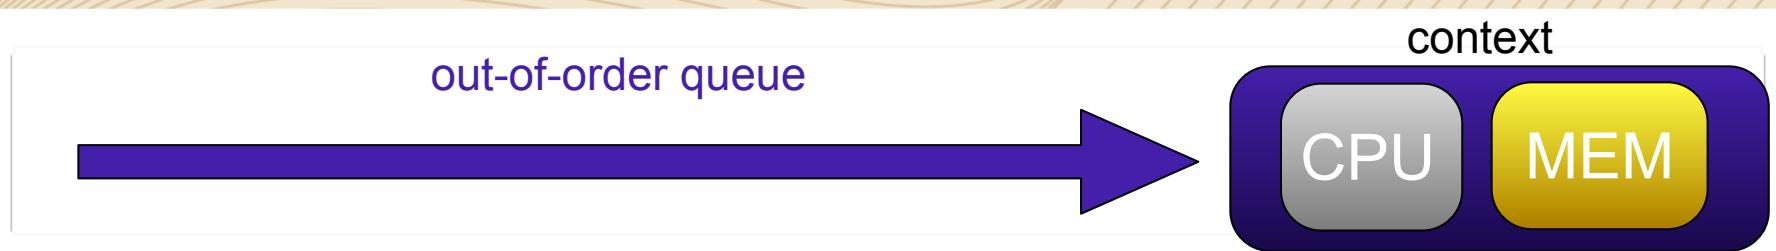
SIGGRAPH ASIA 2009  
the pulse of innovation

# Other Usage Patterns Which Require Synchronisation



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Single Device Out-of-Order Usage Model

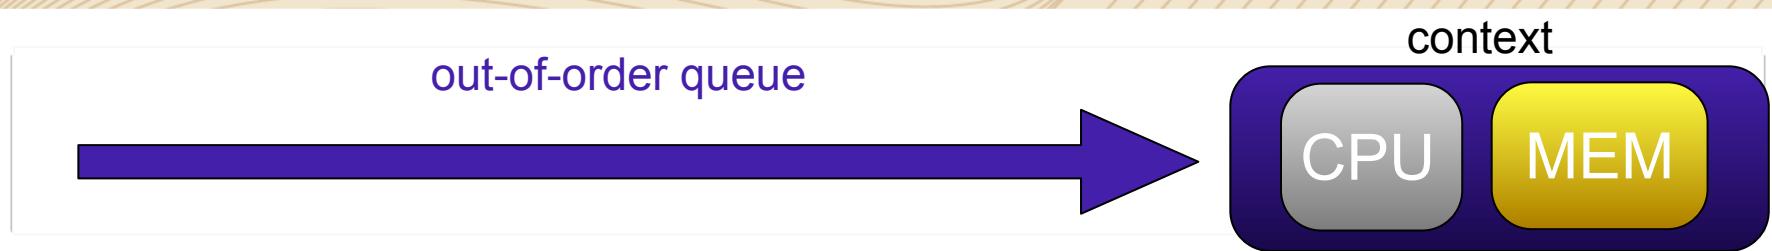


- **1x Out-of-Order Queue, 1x Context, 1x Device**
  - Same as before but with an **out-of-order** queue
  - All commands execute on single device
  - All memory operations occur in single memory pool
  - Execution order has no guarantees....



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Single Device Out-of-Order Usage Model

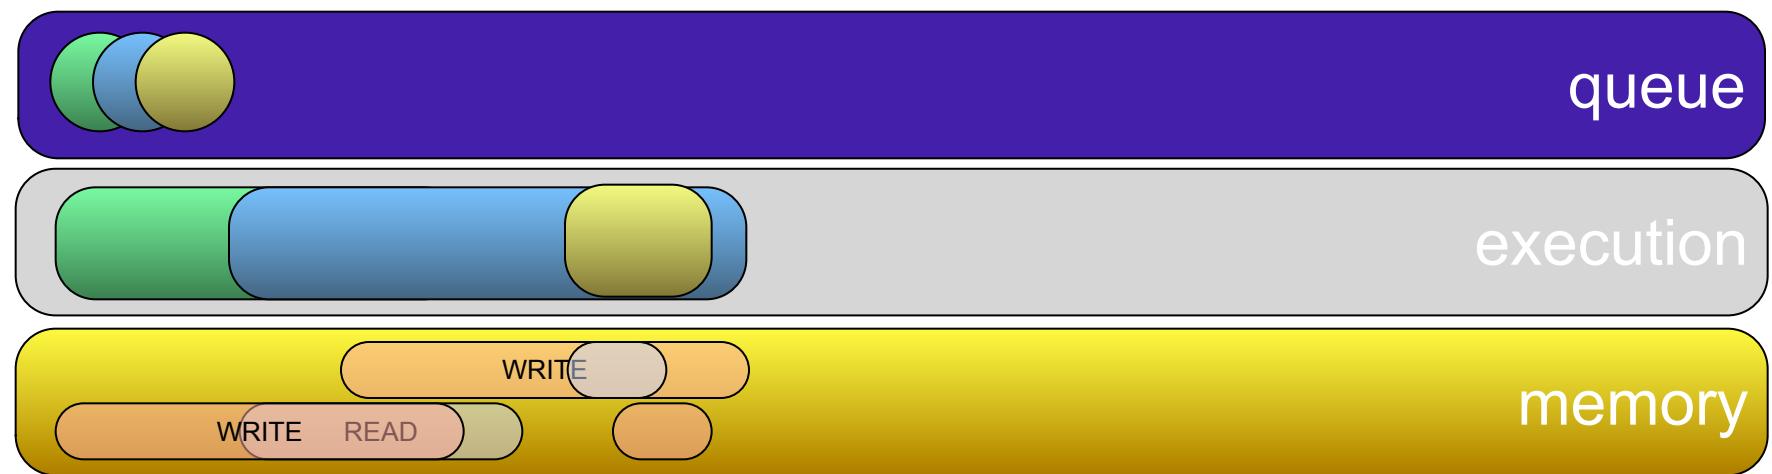


```
cl_uint num_devices;  
cl_device_id devices[1];  
  
err = clGetDeviceIDs(NULL, CL_DEVICE_TYPE_CPU, 1, &devices[0], &num_devices);  
context = clCreateContext(0, 1, devices, NULL, NULL, &err);  
  
cl_command_queue queue_cpu;  
queue_cpu = clCreateCommandQueue(context, devices[0],  
CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE, &err);  
  
/* ... enqueue work ... */
```



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Single Device Out-of-Order Usage Model



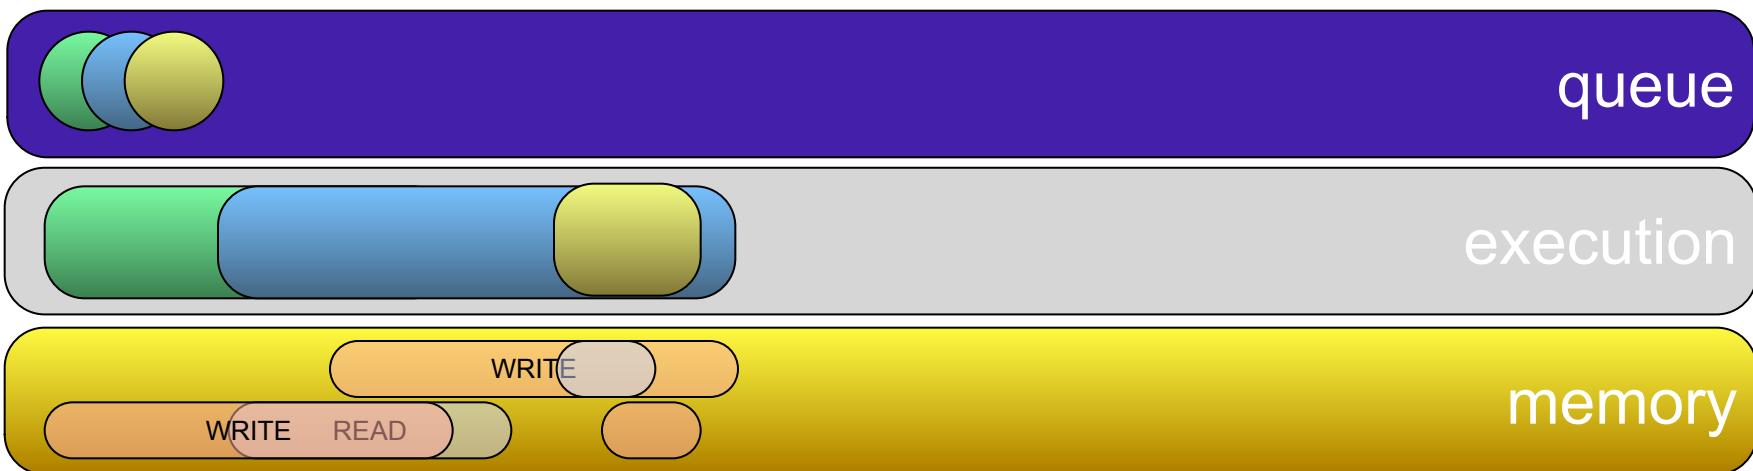
time



SIGGRAPH ASIA 2009  
the pulse of innovation

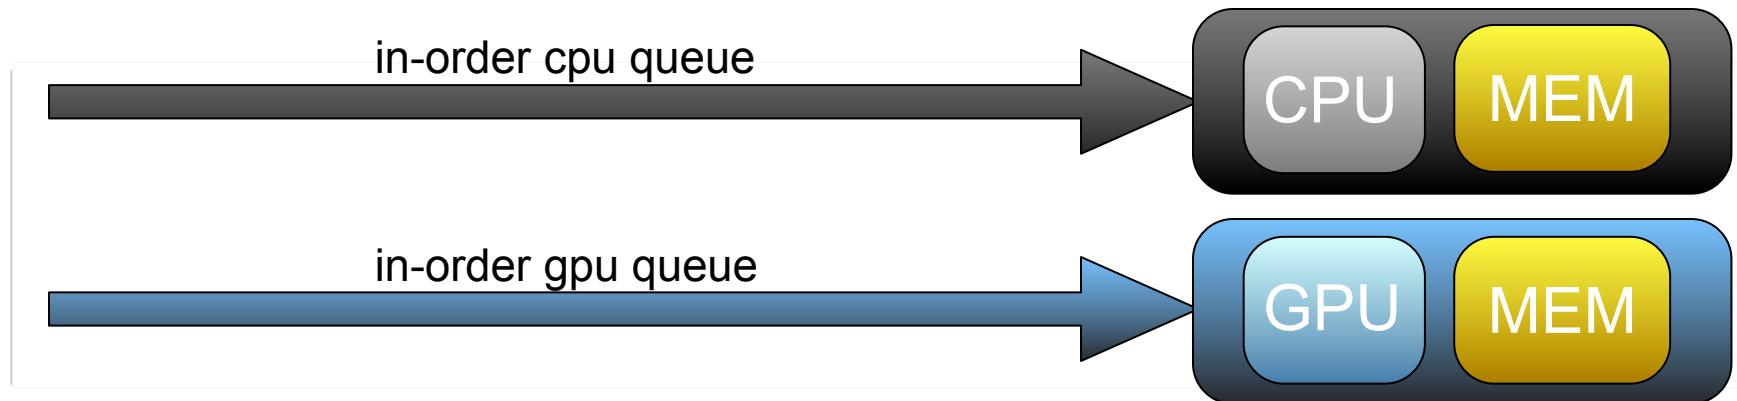
# Single Device Out-of-Order Usage Model

- Device starts executing commands as soon as it can
- Memory transactions overlap and clobber data!



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Separate Multi-Device Usage Model

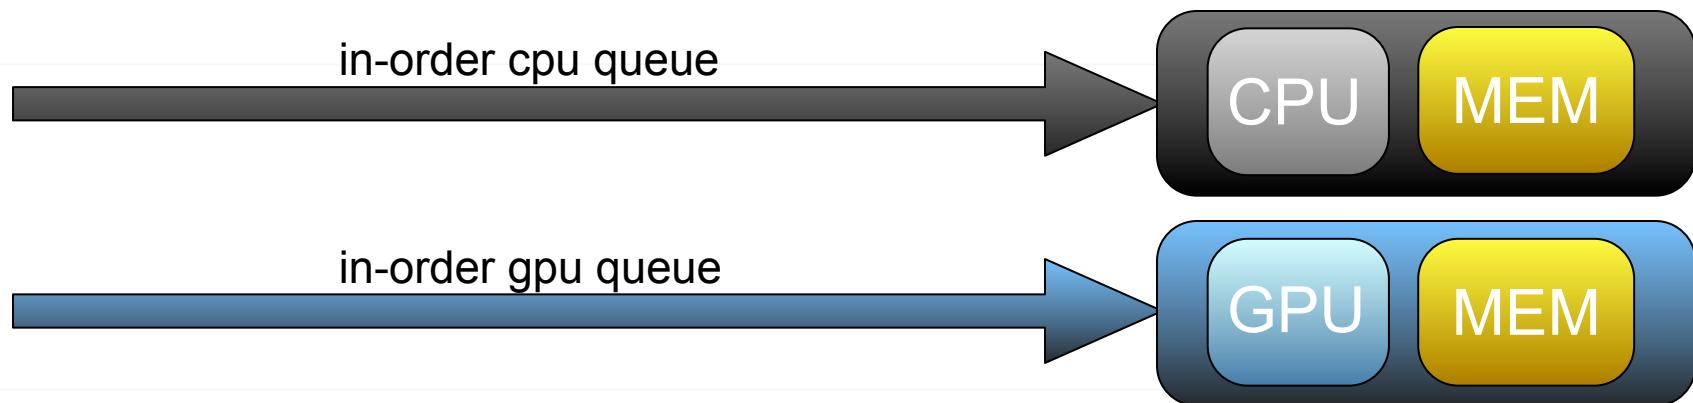


- **2x In-Order Queues, 2x Separate Contexts, 2x Devices**
  - Commands execute on the device associated with the queue
  - Memory operations occur in two **separate memory** pools
  - Completely **separate queues in different contexts**



**SIGGRAPHASIA2009**  
the pulse of innovation

# Separate Multi-Device Usage Model



```
cl_uint num_devices;
cl_device_id devices[2];

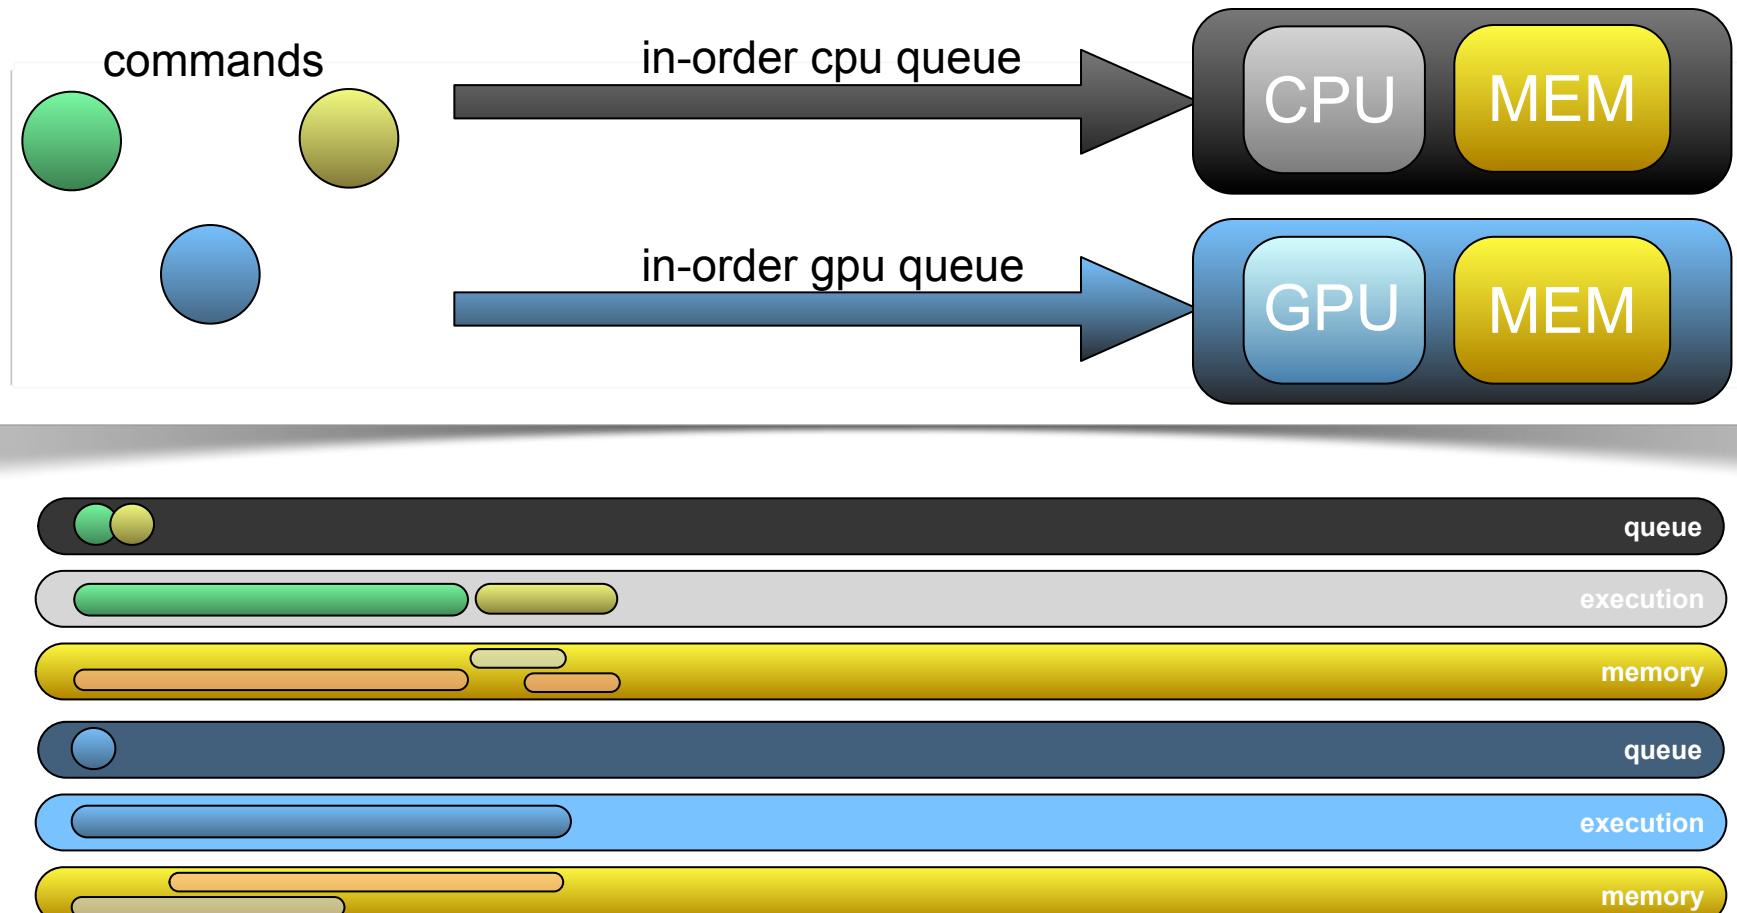
err = clGetDeviceIDs(NULL, CL_DEVICE_TYPE_CPU, 1, &devices[0], &num_devices);
err = clGetDeviceIDs(NULL, CL_DEVICE_TYPE_GPU, 1, &devices[1], &num_devices);
context_cpu = clCreateContext(0, 1, &devices[0], NULL, NULL, &err);
context_gpu = clCreateContext(0, 1, &devices[1], NULL, NULL, &err);

cl_command_queue queue_cpu, queue_gpu;
queue_cpu = clCreateCommandQueue(context_cpu, devices[0], 0 /* IN-ORDER */, &err);
queue_gpu = clCreateCommandQueue(context_gpu, devices[1], 0 /* IN-ORDER */, &err);
```



SIGGRAPH ASIA 2009  
the pulse of innovation

# Separate Multi-Device Usage Model



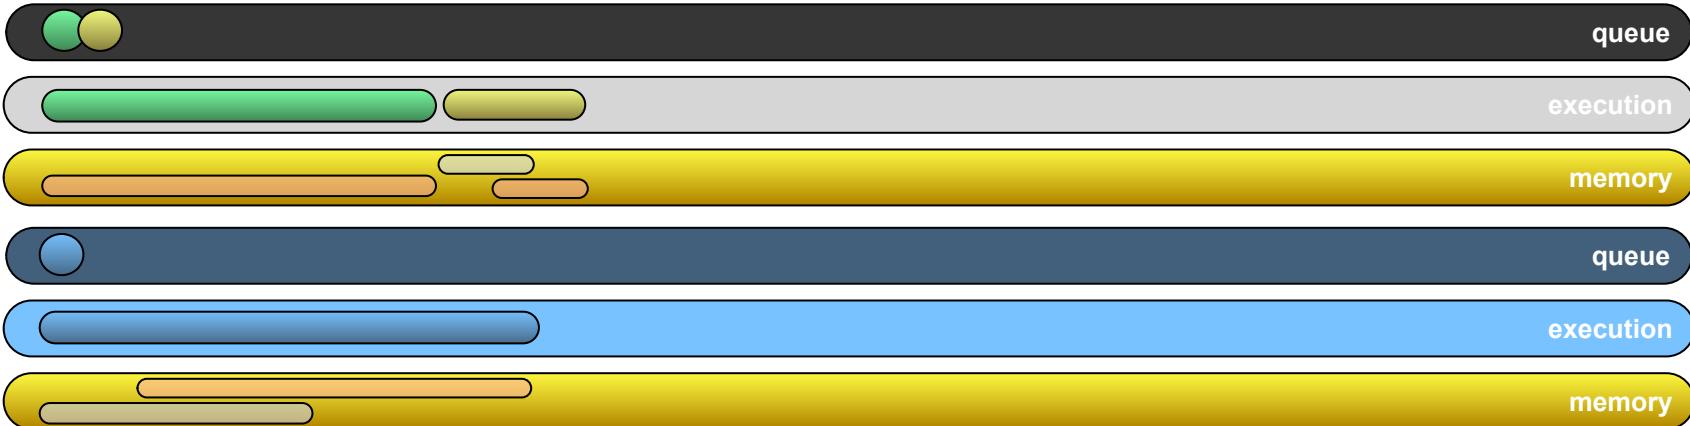
time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Separate Multi-Device Usage Model

- Command queues can't synchronise across contexts!
- Neither device sees the memory pool of the other!
- Won't work unless you use `clFinish()`, and copy across contexts!
- In short, this is not what you want! Don't do this!

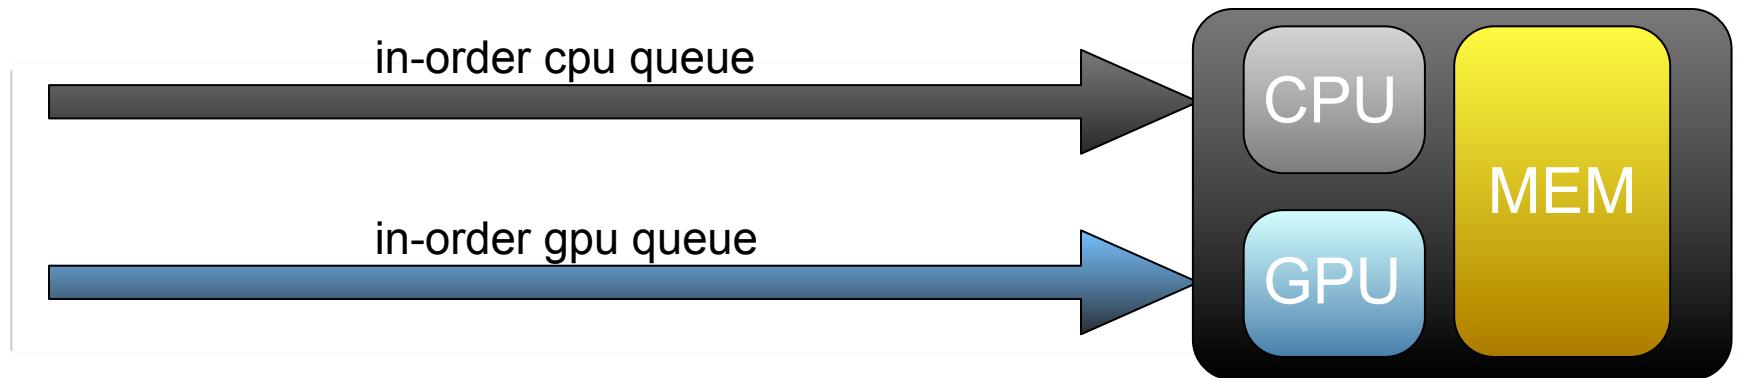


time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Cooperative Multi-Device Usage Model

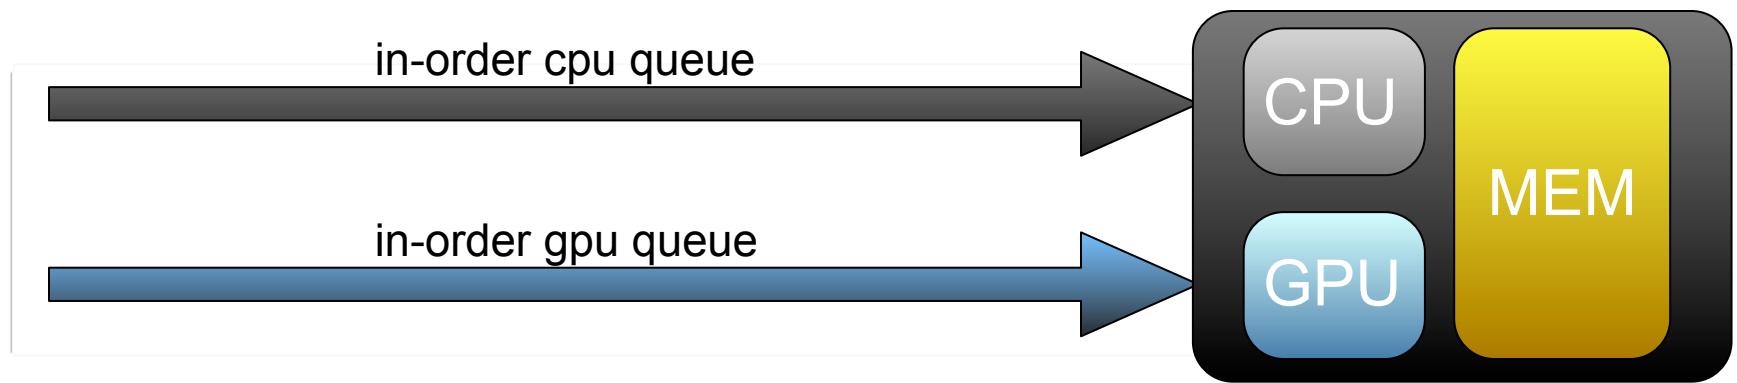


- **2x In-Order Queues, 1x Combined Context, 2x Devices**
  - Commands execute on the device associated with the queue
  - Memory operations occur in one **combined memory** pool
  - **Combined memory** pool being modified by multiple devices



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Cooperative Multi-Device Usage Model



```
cl_uint num_devices;
cl_device_id devices[2];

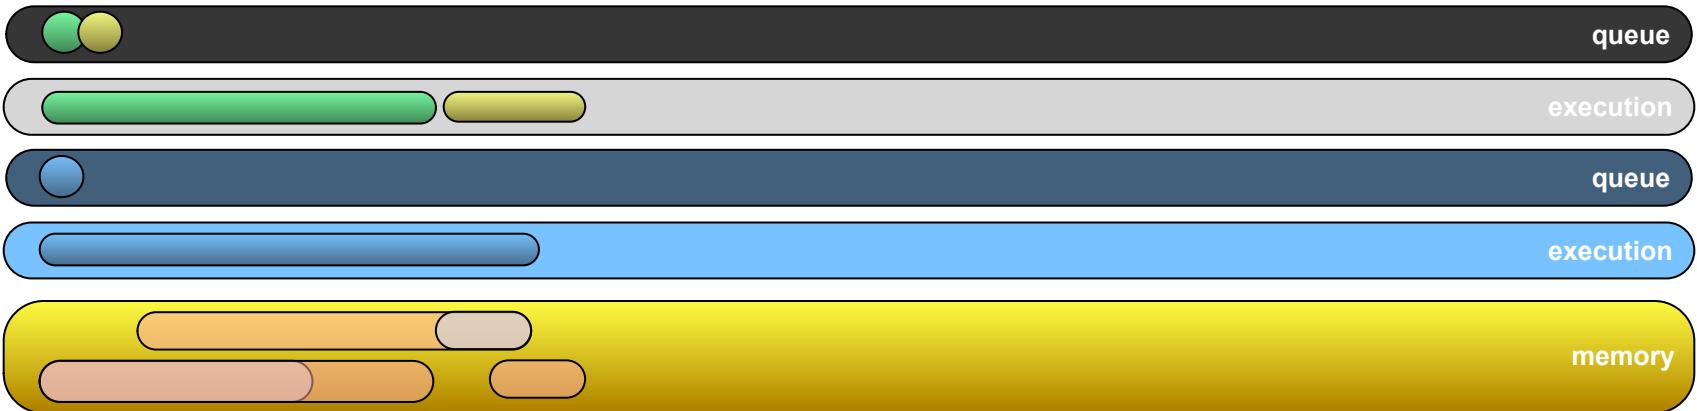
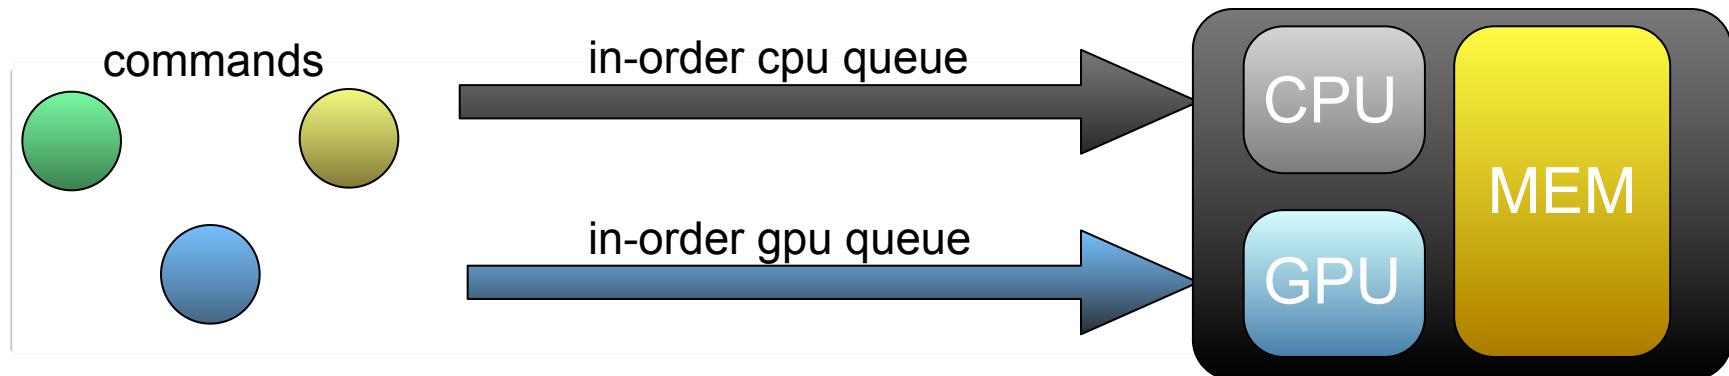
err = clGetDeviceIDs(NULL, CL_DEVICE_TYPE_CPU, 1, &devices[0], &num_devices);
err = clGetDeviceIDs(NULL, CL_DEVICE_TYPE_GPU, 1, &devices[1], &num_devices);
context = clCreateContext(0, 2, devices, NULL, NULL, &err);

cl_command_queue queue_cpu, queue_gpu;
queue_cpu = clCreateCommandQueue(context, devices[0], 0 /* IN-ORDER */, &err);
queue_gpu = clCreateCommandQueue(context, devices[1], 0 /* IN-ORDER */, &err);
```



SIGGRAPH ASIA 2009  
the pulse of innovation

# Cooperative Multi-Device Usage Model



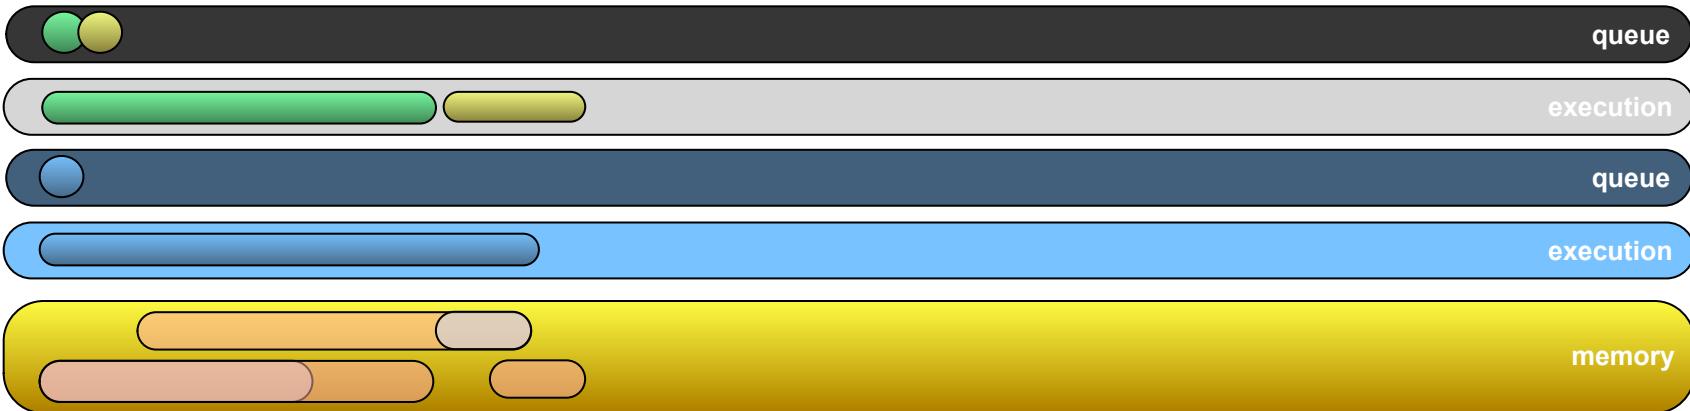
time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Cooperative Multi-Device Usage Model

- Still wrong!
- Both devices start executing commands as soon as they can!
- Memory transactions overlap and clobber combined memory!



time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Synchronisation Mechanisms



**SIGGRAPHASIA2009**  
*the pulse of innovation*

# Synchronisation Mechanisms

- **Command Queue Control Methods**

- **clFlush()**

Send all commands in the queue to the compute device

- **clFinish()**

Flush, and then wait for all commands to finish



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Synchronisation Mechanisms

- **Command Queue Control Methods**

- `clFlush()`

- Send all commands in the queue to the compute device

- `clFinish()`

- Flush, and then wait for all commands to finish

- **Both are brute force, and lack fine-grained control**



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Synchronisation Mechanisms

- **Command Execution Barriers**

- `cIEnqueueBarrier()`

Enqueue a fence which insures that all preceding commands in the queue are complete before any commands which get enqueued afterwards are processed



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Synchronisation Mechanisms

- **Event Based Synchronisation**

- **Event objects**

Unique objects which can be used to determine command status

- **Event wait lists**

Array of events used to indicate commands which must be complete before further commands are allowed to proceed



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- Event objects are used to determine command status

- **clGetEventStatus()**

Returns command status for an event

**CL\_QUEUED**

- Command is in a queue

**CL\_SUBMITTED**

- Command has been submitted to device

**CL\_RUNNING**

- Device is currently executing command

**CL\_COMPLETE**

- Command has finished execution



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods can return event objects

- `clEnqueueNDRangeKernel`
- `clEnqueueTask`
- `clEnqueueNativeKernel`
- `clEnqueueCopyImageToBuffer`
- `clEnqueueCopyBufferToImage`
- `clEnqueueRead {Image|Buffer}`
- `clEnqueueWrite {Image|Buffer}`
- `clEnqueueMap {Image|Buffer}`
- `clEnqueueCopy {Image|Buffer}`
- `clEnqueueCopyImageToBuffer`
- `clEnqueueCopyBufferToImage`



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods can return event objects

```
cl_event returned_event;  
  
err = clEnqueueReadBuffer(queue, buffer,  
                           CL_FALSE /* non-blocking */,  
                           0, 0, ptr, , 0,  
                           &returned_event);
```



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods can return event objects

```
cl_event returned_event;  
  
err = clEnqueueReadBuffer(queue, buffer,  
                           CL_FALSE /* non-blocking */,  
                           0, 0, ptr, , 0,  
                           &returned_event);
```

- When blocking is false, call returns immediately!



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- Event objects can be used as synchronisation points
  - `c|WaitForEvents(num_events, event_list)`  
Waits and blocks until all commands identified by events in the given event list are complete



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- Event objects can be used as synchronisation points

```
cl_event read_event;  
  
err = clEnqueueReadBuffer(queue, buffer,  
                           CL_FALSE /* blocking_read */,  
                           0, 0, ptr, 0, 0,  
                           &read_event);  
  
err = clWaitOnEvents(1, &read_event);
```

- Above example is equivalent to blocking call



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods also accept **event wait lists**

```
cl_uint num_events_in_waitlist = 2;  
cl_event event_waitlist[2] = { event_one, event_two };  
  
err = clEnqueueReadBuffer(queue, buffer,  
                          CL_FALSE /* non-blocking */,  
                          0, 0,  
                          num_events_in_waitlist,  
                          event_waitlist, NULL);
```



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods also accept **event wait lists**

```
cl_uint num_events_in_waitlist = 2;  
cl_event event_waitlist[2] = { event_one, event_two };  
  
err = clEnqueueReadBuffer(queue, buffer,  
                          CL_FALSE /* non-blocking */,  
                          0, 0,  
                          num_events_in_waitlist,  
                          event_waitlist, NULL);
```

- Read won't occur until events in wait list are complete!



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods also accept **event wait lists**

```
cl_uint num_events_in_waitlist = 2;
cl_event event_waitlist[2];

err = clEnqueueReadBuffer(queue, buffer0, CL_FALSE /* non-blocking */, 0, 0,
                         0, NULL, &event_waitlist[0]);

err = clEnqueueReadBuffer(queue, buffer1, CL_FALSE /* non-blocking */, 0, 0,
                         0, NULL, &event_waitlist[1]);

/* last read buffer waits on previous two read buffer events */
err = clEnqueueReadBuffer(queue, buffer2, CL_FALSE /* non-blocking */, 0, 0,
                         num_events_in_waitlist, event_waitlist, NULL);
```



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- All `clEnqueue()` methods also accept **event wait lists**

```
cl_uint num_events_in_waitlist = 2;
cl_event event_waitlist[2];

err = clEnqueueReadBuffer(queue, buffer0, CL_FALSE /* non-blocking */, 0, 0,
                         0, NULL, &event_waitlist[0]);

err = clEnqueueReadBuffer(queue, buffer1, CL_FALSE /* non-blocking */, 0, 0,
                         0, NULL, &event_waitlist[1]);

/* last read buffer waits on previous two read buffer events */
err = clEnqueueReadBuffer(queue, buffer2, CL_FALSE /* non-blocking */, 0, 0,
                         num_events_in_waitlist, event_waitlist, NULL);
```

- Last read buffer waits on previous other two to complete!



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- Events can also provide profiling information

```
cl_event read_event;
cl_ulong start, end;

queue = clCreateCommandQueue(context, device, CL_QUEUE_PROFILING_ENABLED, ...
err = clEnqueueReadBuffer(queue, buffer, 0, 0, 0, 0, NULL, &read_event);
err = clWaitOnEvents(1, read_event);

err = clGetEventProfilingInfo(read_event, CL_PROFILING_COMMAND_END,
                           sizeof(cl_ulong), &end, NULL);

err = clGetEventProfilingInfo(read_event, CL_PROFILING_COMMAND_START,
                           sizeof(cl_ulong), &start, NULL);

float milliseconds = (end - start) * 1.0e-6f;
```



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- **Event-Based Barrier Methods**

- **clEnqueueWaitList()**

Enqueue a list of events to wait for such that all events need to complete before this particular command can be executed

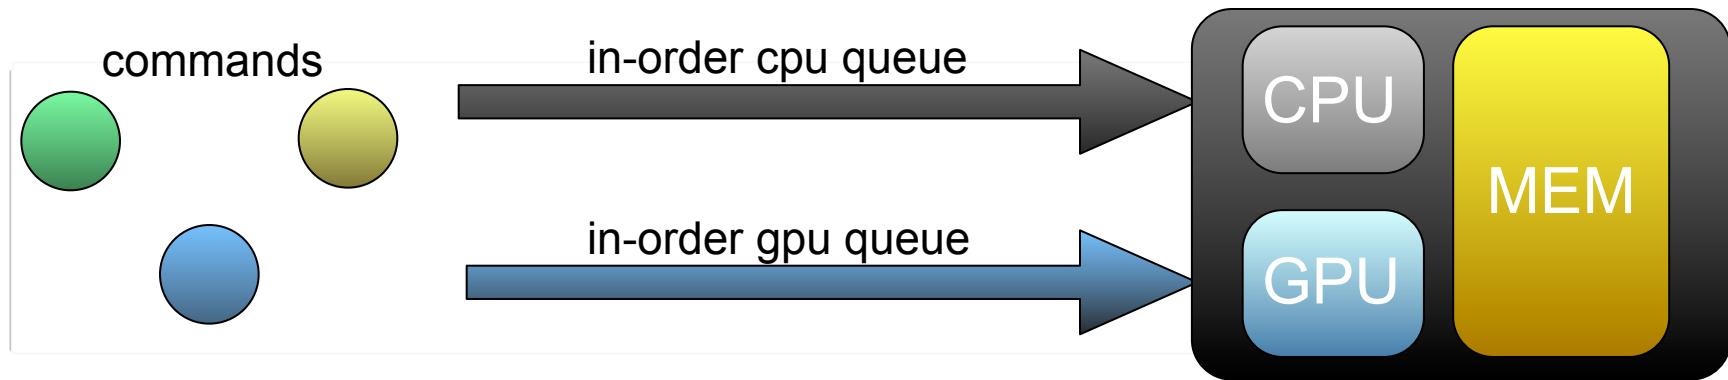
- **clEnqueueMarker()**

Enqueue a command which marks this location in the queue with a unique event object that can be used for synchronisation



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Cooperative Multi-Device Usage Model



```
cl_event event0_cpu, event1_gpu;  
  
err = clEnqueueNDRangeKernel(queue_cpu, kernel0_cpu, 2, NULL, global, local,  
    0, NULL, &event0_cpu);  
  
err = clEnqueueNDRangeKernel(queue_gpu, kernel1_gpu, 2, NULL, global, local,  
    1, &event0_cpu, &event1_gpu);  
  
err = clEnqueueNDRangeKernel(queue_cpu, kernel2_cpu, 2, NULL, global, local,  
    1, &event1_gpu, NULL);  
  
clFlush(queue_cpu); clFlush(queue_gpu);
```

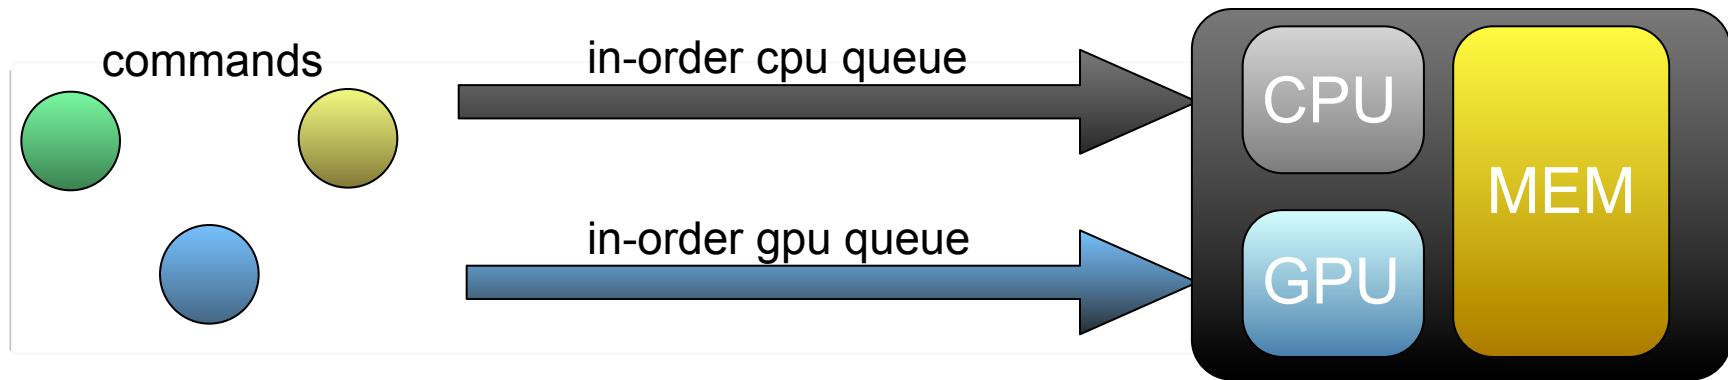


time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Cooperative Multi-Device Usage Model



```
cl_event event0_cpu, event1_gpu;  
  
err = clEnqueueNDRangeKernel(queue_cpu, kernel0_cpu, 2, NULL, global, local,  
    0, NULL, &event0_cpu);  
  
err = clEnqueueNDRangeKernel(queue_gpu, kernel1_gpu, 2, NULL, global, local,  
    1, &event0_cpu, &event1_gpu);  
  
err = clEnqueueNDRangeKernel(queue_cpu, kernel2_cpu, 2, NULL, global, local,  
    1, &event1_gpu, NULL);  
  
clFlush(queue_cpu); clFlush(queue_gpu);
```



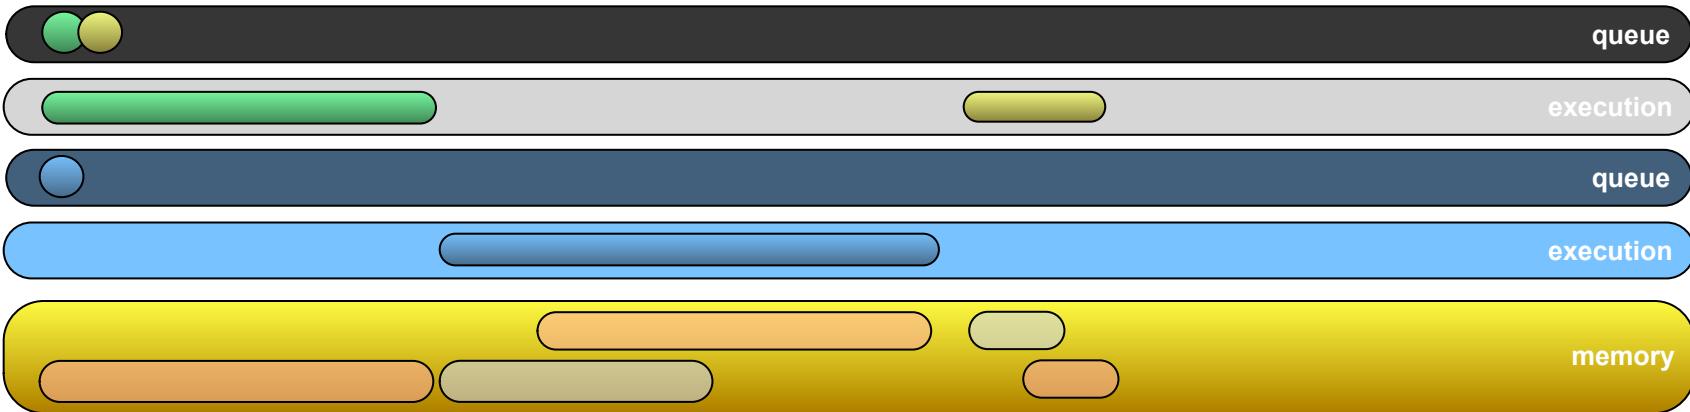
time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Cooperative Multi-Device Usage Model

- Event wait lists provide synchronised execution
- Execution is dependant on prior command execution status
- View of shared memory pool is consistent during execution



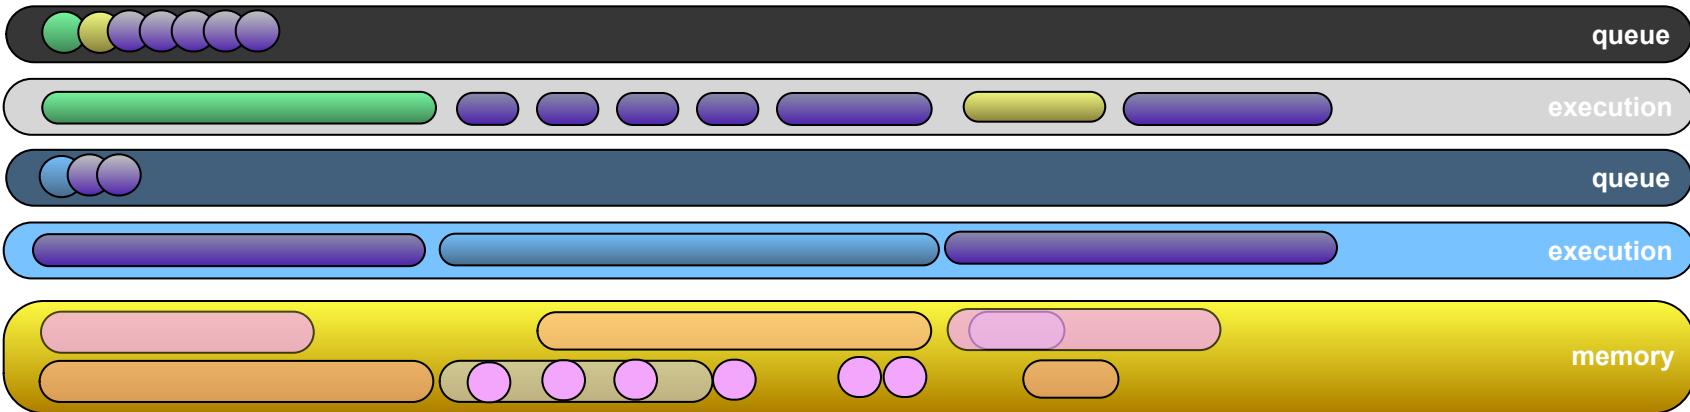
time



**SIGGRAPHASIA2009**  
the pulse of innovation

# Cooperative Multi-Device Usage Model

- Event wait lists provide synchronised execution
- Execution is dependant on prior command execution status
- View of shared memory pool is consistent during execution
- Could submit other work to fill in the time gaps!



time



SIGGRAPH ASIA 2009  
the pulse of innovation

# Event-Based Synchronisation

- **Event Model Summary**

- Event objects identify unique commands in a queue
- Event objects can be used to determine command status
- Event objects also provide profiling information
- Event objects can be used as synchronisation points



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Event-Based Synchronisation

- **Event Model Summary**

- Event objects identify unique commands in a queue
- Event objects can be used to determine command status
- Event objects also provide profiling information
- Event objects can be used as synchronisation points
- **Event objects provide command-level control!**
- **Event wait lists allow execution graphs!**



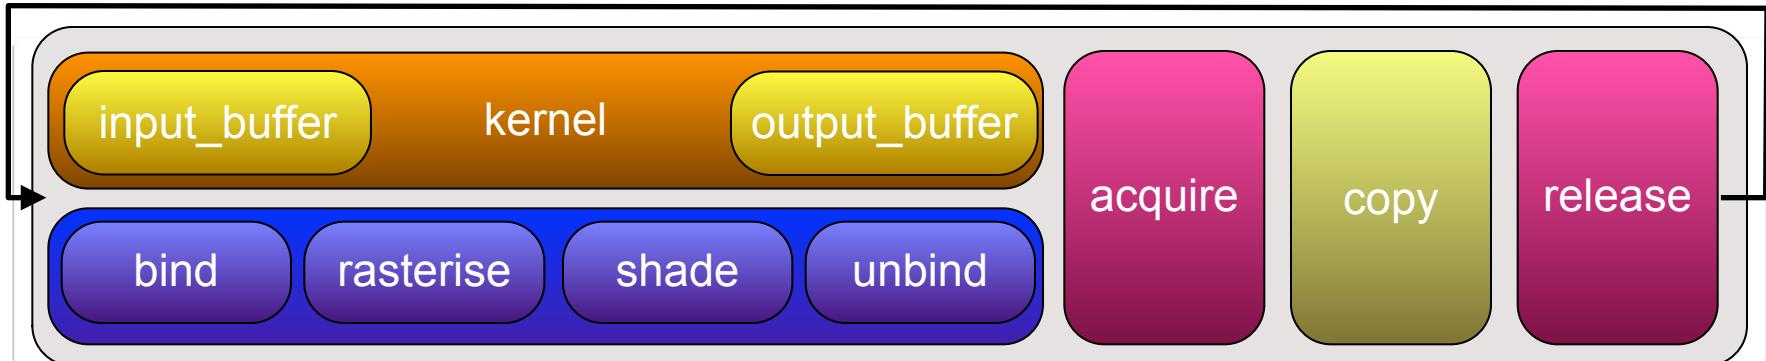
**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Example Usage



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Format Conversion for OpenGL Interop

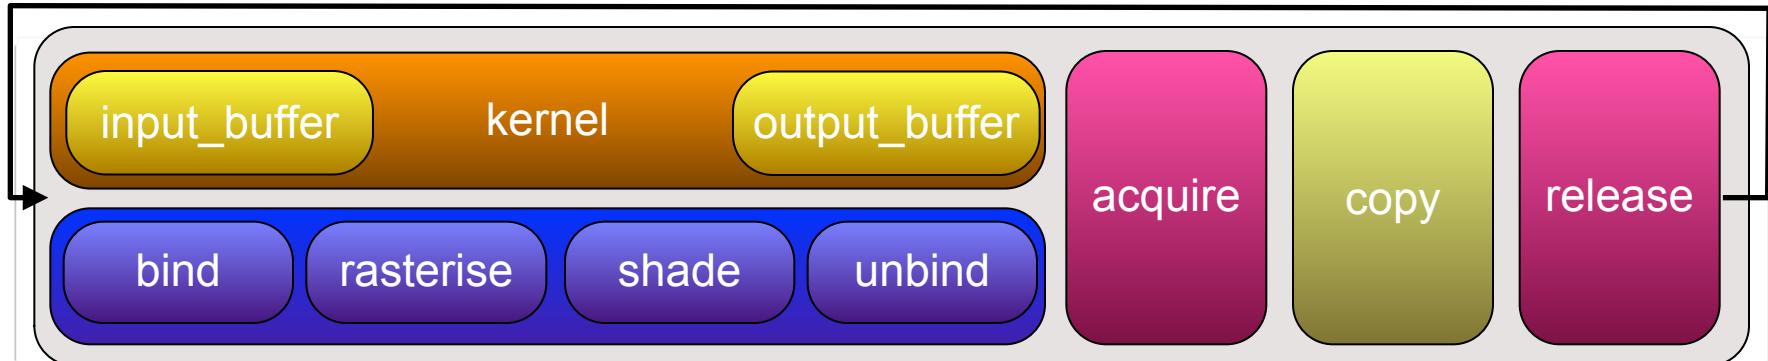


- Some **OpenCL** kernels may be better off using buffers
- Some **OpenGL** commands may be better off using textures
- Use **acquire** and **release** interop methods to hint API usage
- Use **copy** format methods to convert memory



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Format Conversion for OpenGL Interop



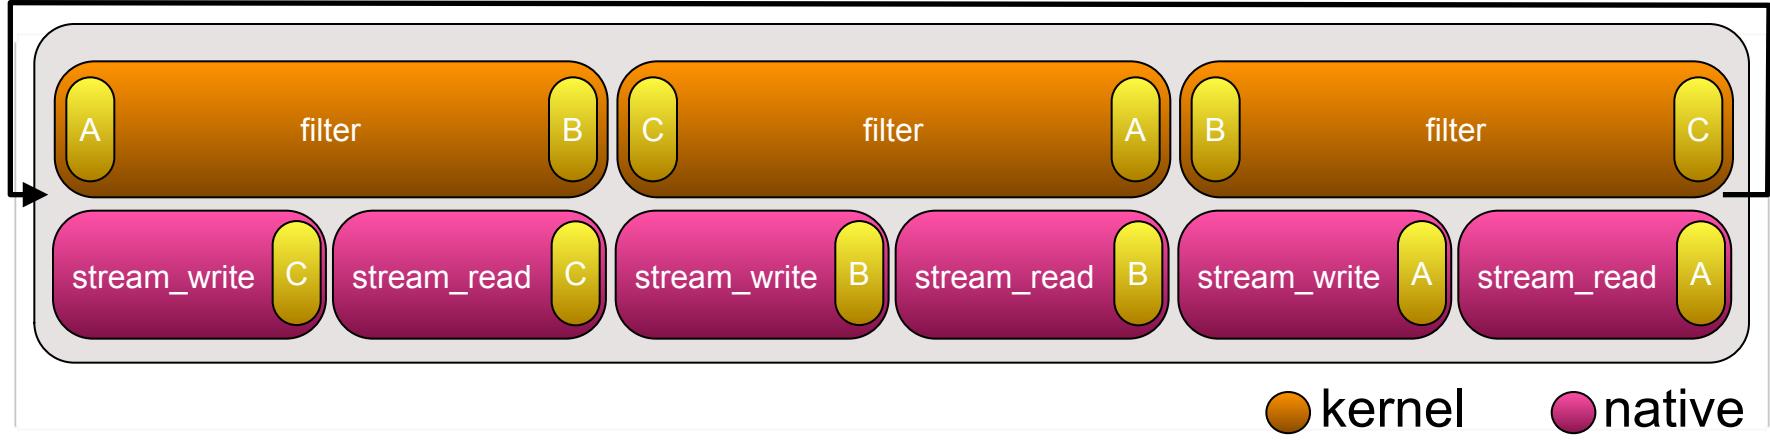
```
size_t origin[] = { 0, 0, 0 };
size_t region[] = { width, height, 1 };
cl_event convert_event;

err = clEnqueueAcquireGLObjects(queue, 1, &image, 0, 0, 0);
err = clEnqueueCopyBufferToImage(queue, buffer, image, 0,
                                 origin, region, 0, NULL, &convert_event);
err = clEnqueueReleaseGLObjects(queue, 1, &image, 0, 0, 0);
```



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data

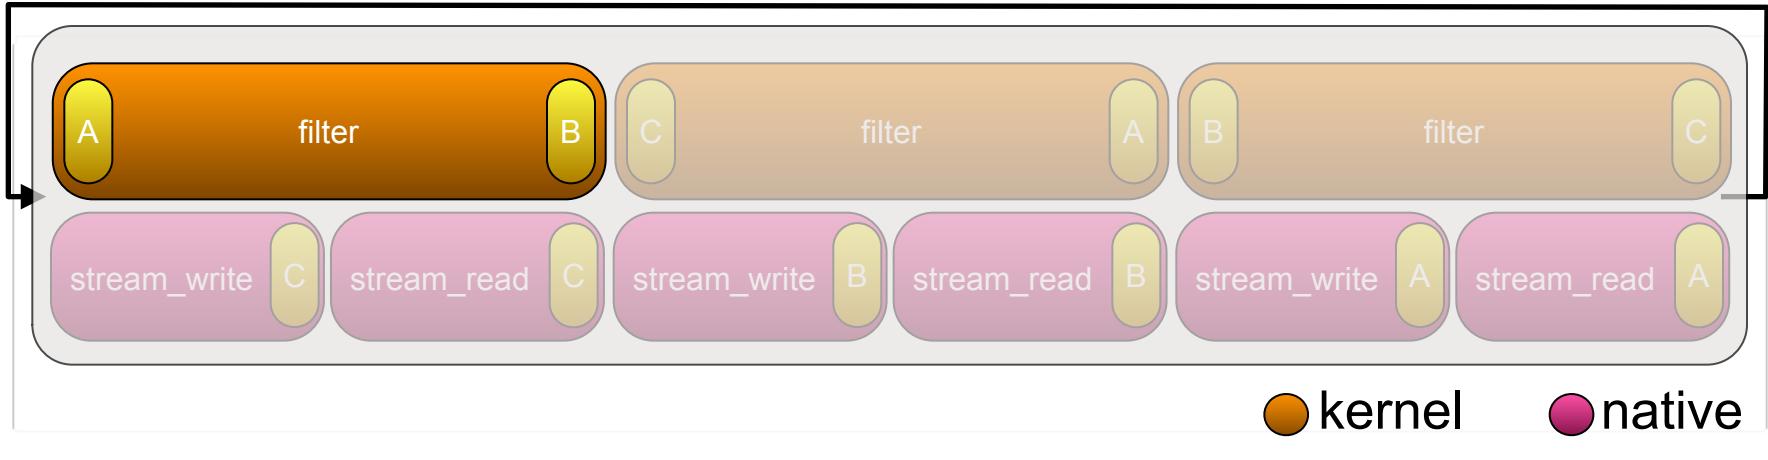


- Overlap streaming data operations with compute
- Use device to do all computation with **kernel**s
- Use host to do all streaming operations w/**native** functions
- Use **event wait lists** for execution order



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data

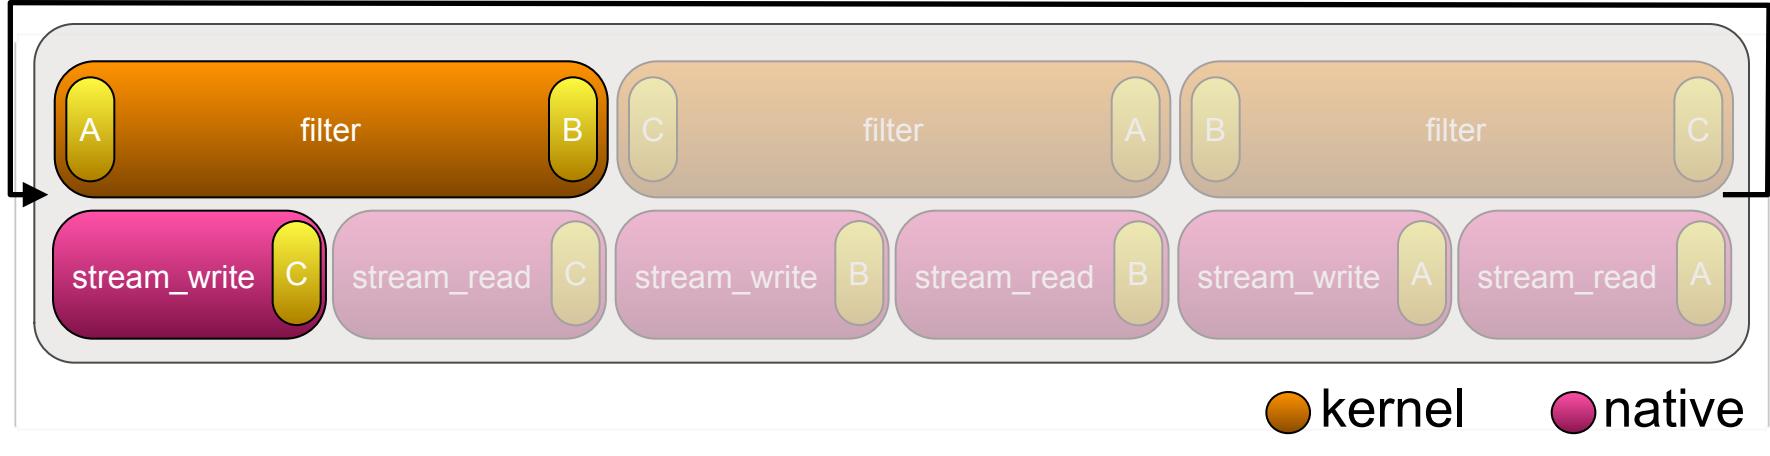


- On the device:  
**filter kernel** reads from **cl\_mem A** writes to **cl\_mem B**



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Streaming Data

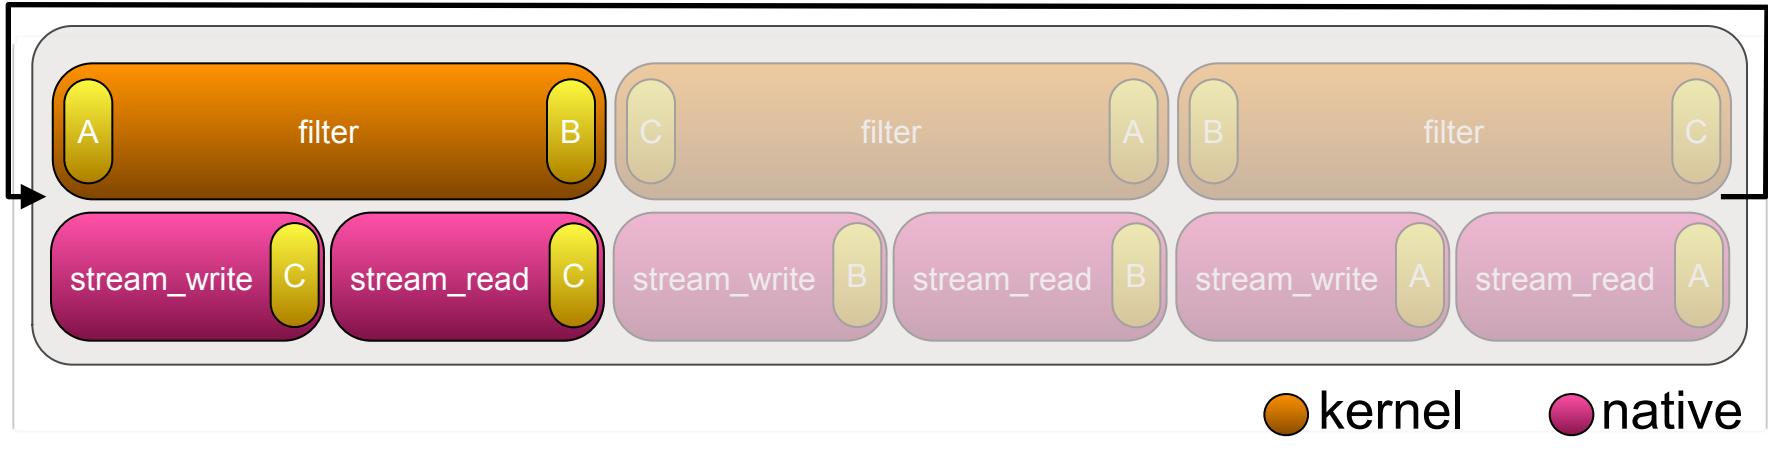


- On the device:  
**filter kernel** reads from **cl\_mem A** writes to **cl\_mem B**
- Meanwhile, on the host:  
**stream\_write** method outputs previous **cl\_mem C**



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data

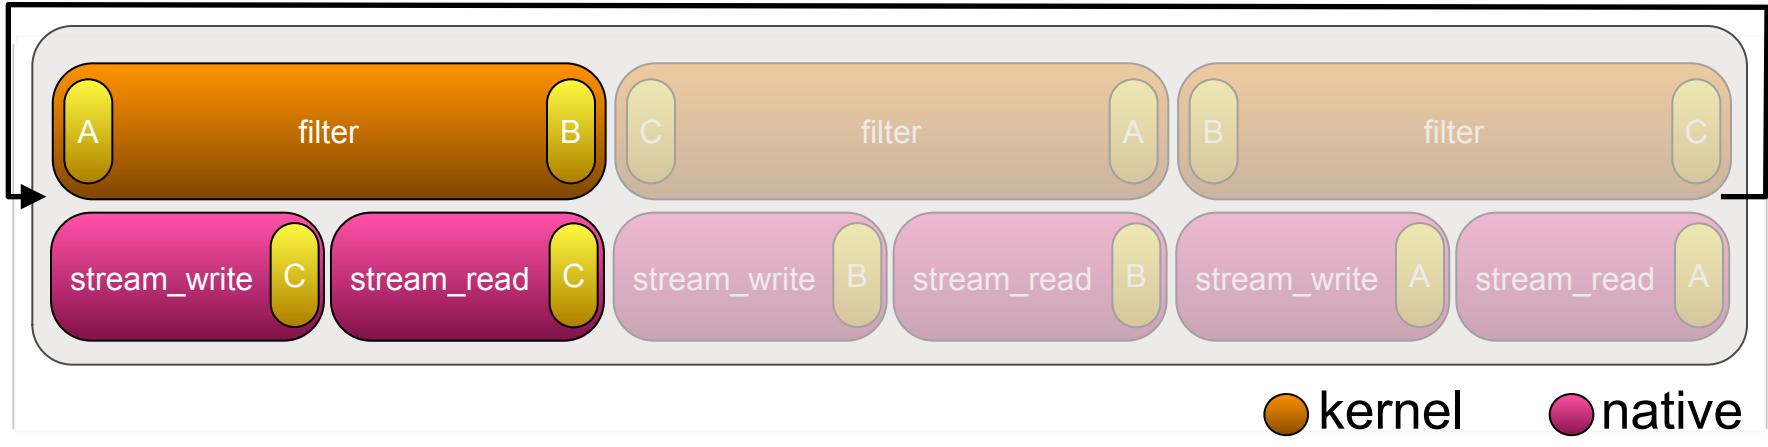


- On the device:  
**filter kernel** reads from **cl\_mem A** writes to **cl\_mem B**
- Meanwhile, on the host:  
**stream\_write** method outputs previous **cl\_mem C**  
**stream\_read** method prefetches and fills **cl\_mem C**



SIGGRAPH ASIA 2009  
the pulse of innovation

# Streaming Data



- Continue and alternative memory locations....

**filter kernel:** ( A - B )

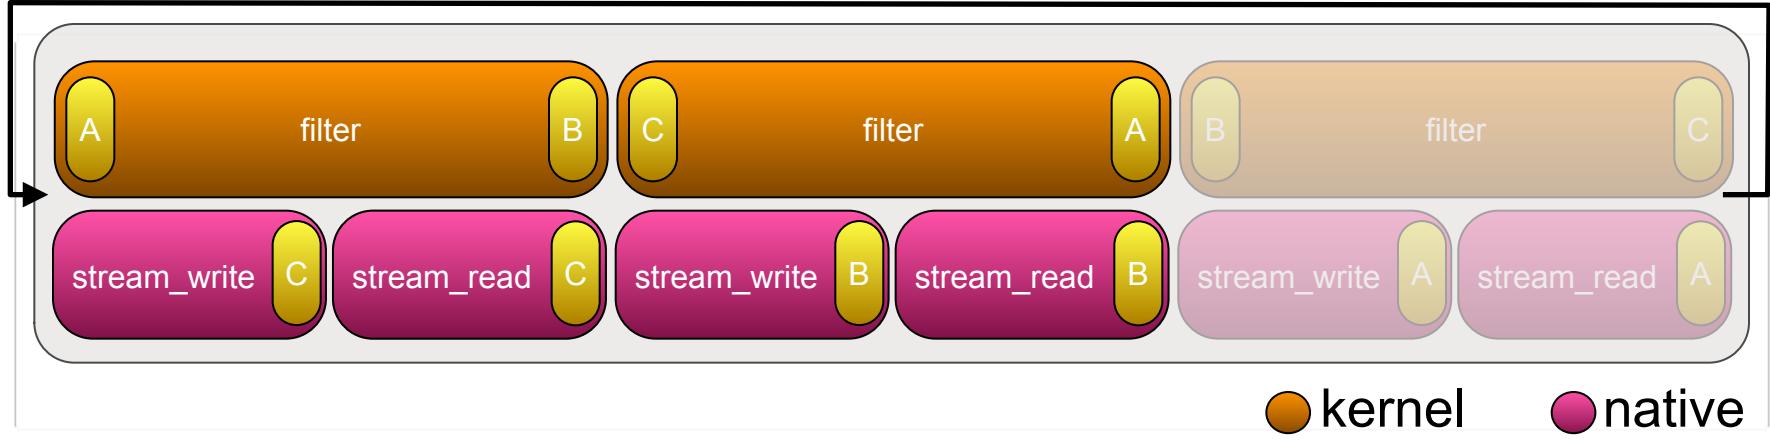
**stream\_write:** ( C )

**stream\_read:** ( C )



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data



- Continue and alternative memory locations....

**filter kernel:** ( A - B ) ( C - A )

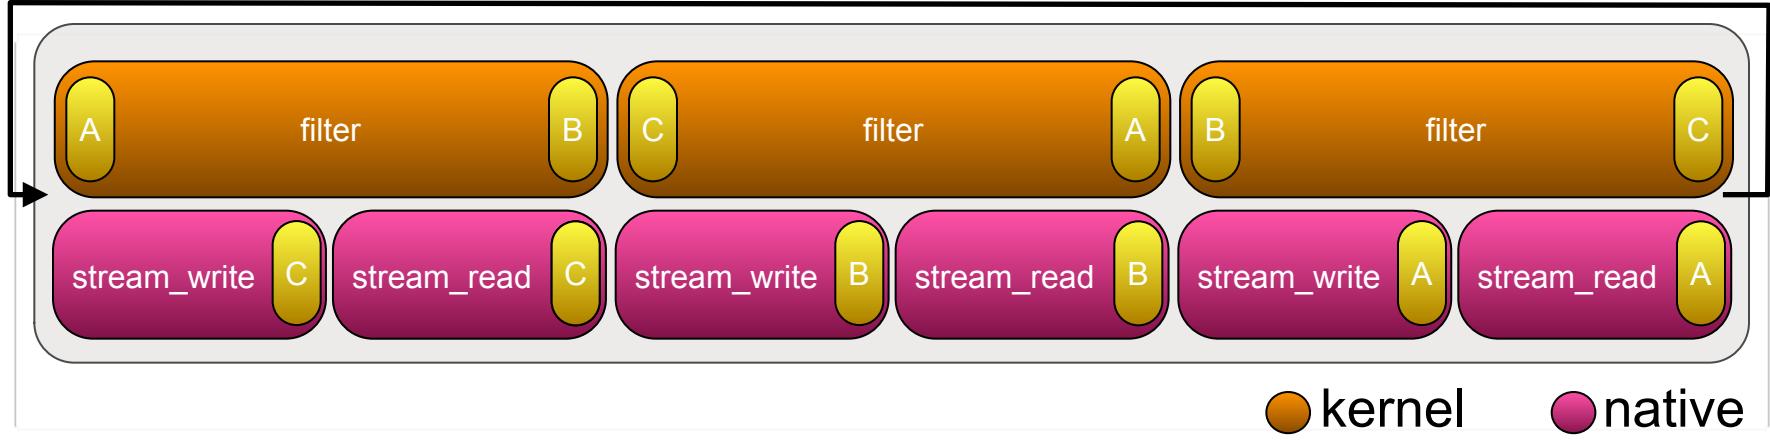
**stream\_write:** ( C ) ( B )

**stream\_read:** ( C ) ( B )



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Streaming Data



- Continue and alternative memory locations....

**filter kernel:** ( A - B ) ( C - A ) ( B - C )

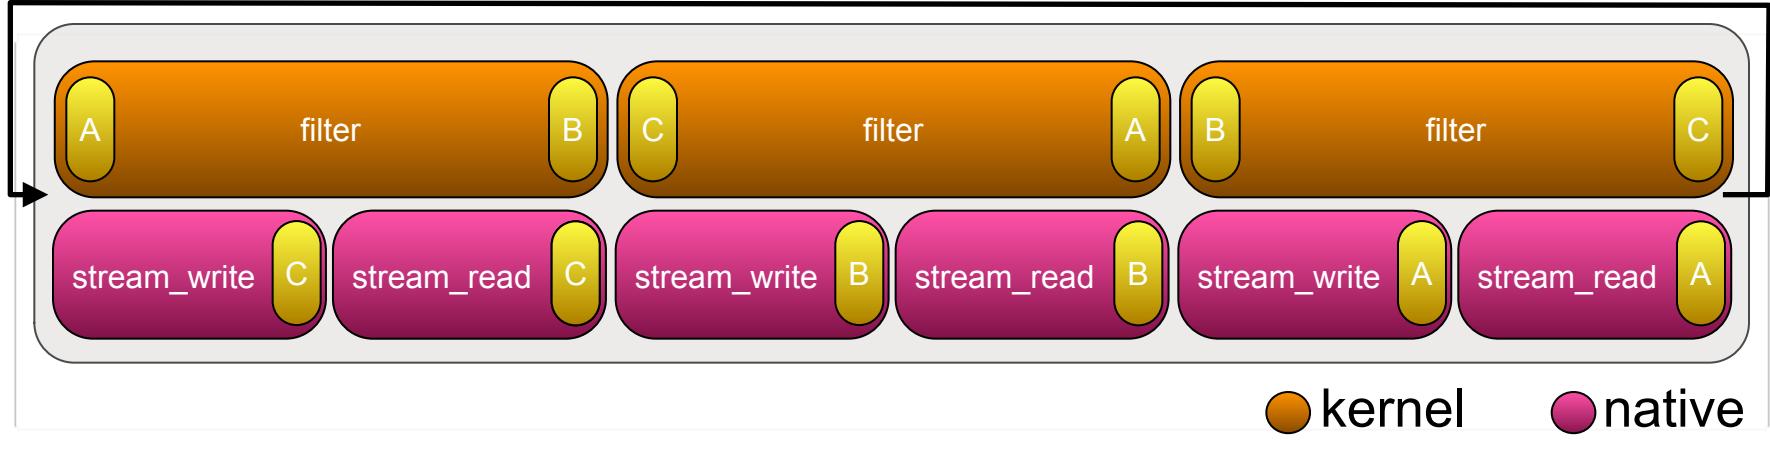
**stream\_write:** ( C ) ( B ) ( A )

**stream\_read:** ( C ) ( B ) ( A )



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data

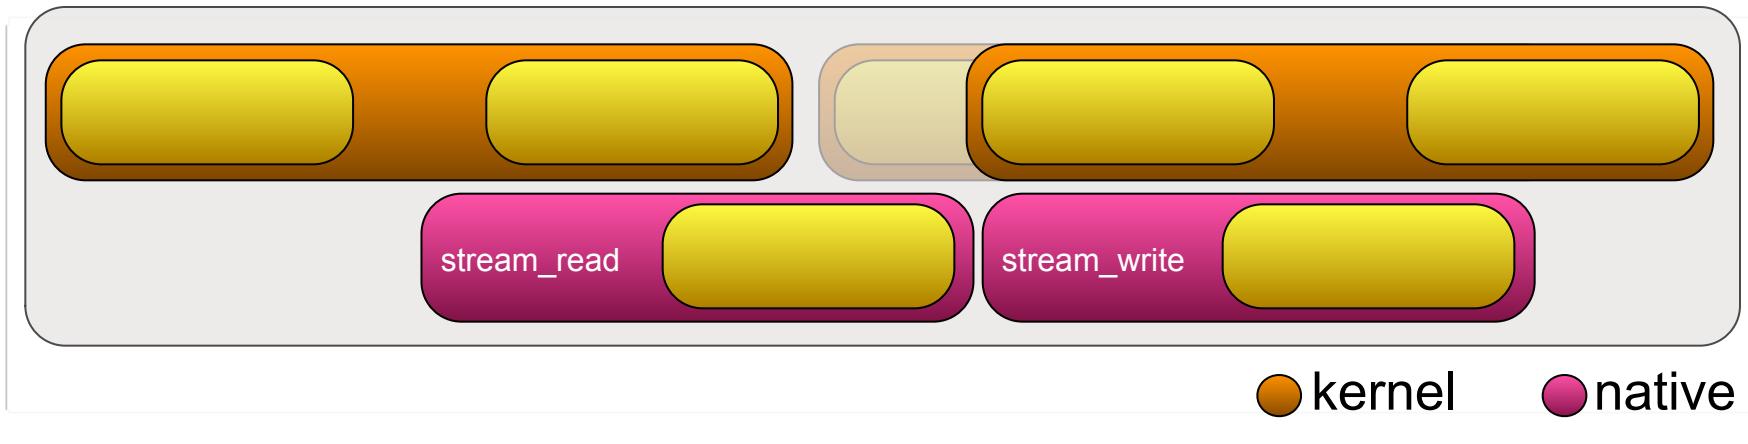


- **Some caveats....**
  - Size of data stream must be large to offset transfer costs
  - Complexity of filter must be large to offset launch costs
  - Both **stream\_write** and **stream\_read** methods must be timely



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data

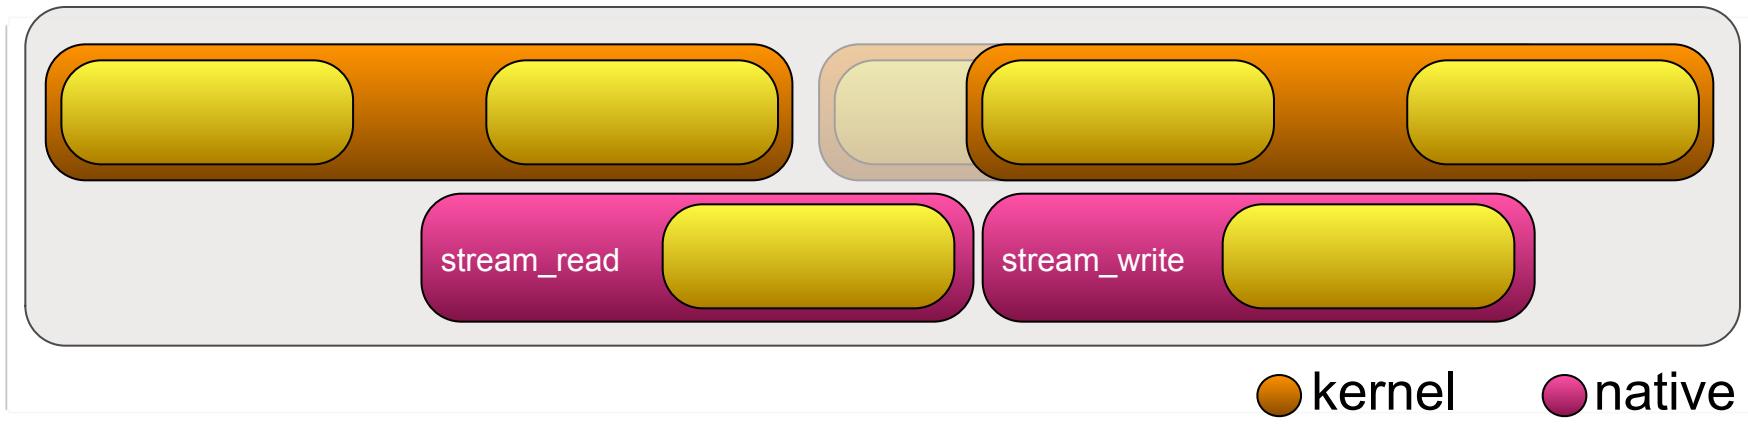


- If data size is too large....
  - `stream_read` can't fill full data size in memory
  - `stream_write` can't output full data size



**SIGGRAPH ASIA 2009**  
the pulse of innovation

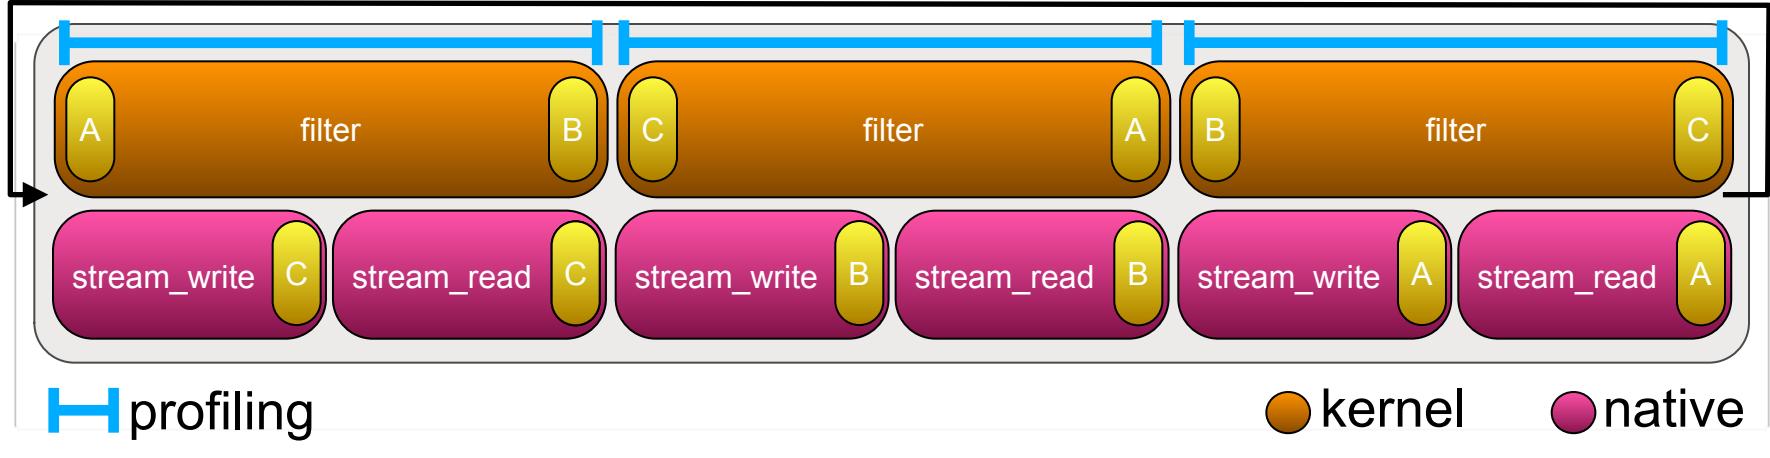
# Streaming Data



- **Don't want to delay device kernel**
  - Not willing to sacrifice efficiency
  - Might have strict time constraints for device results



# Streaming Data

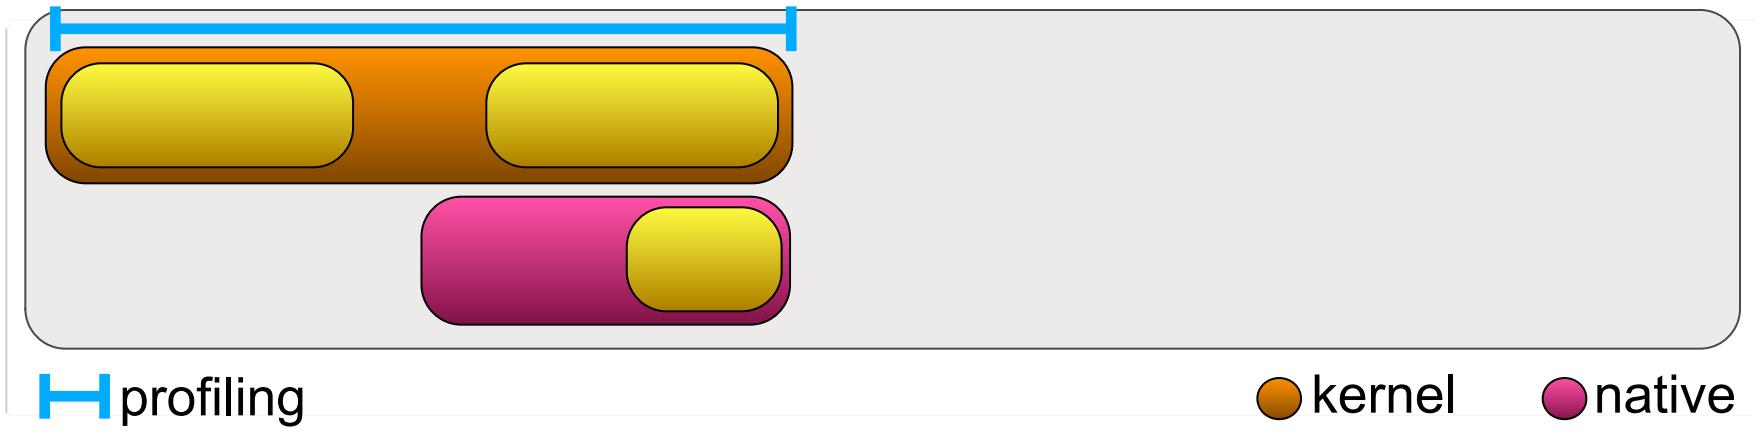


- Use **event** profiling to identify time window for **stream operations**



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# Streaming Data

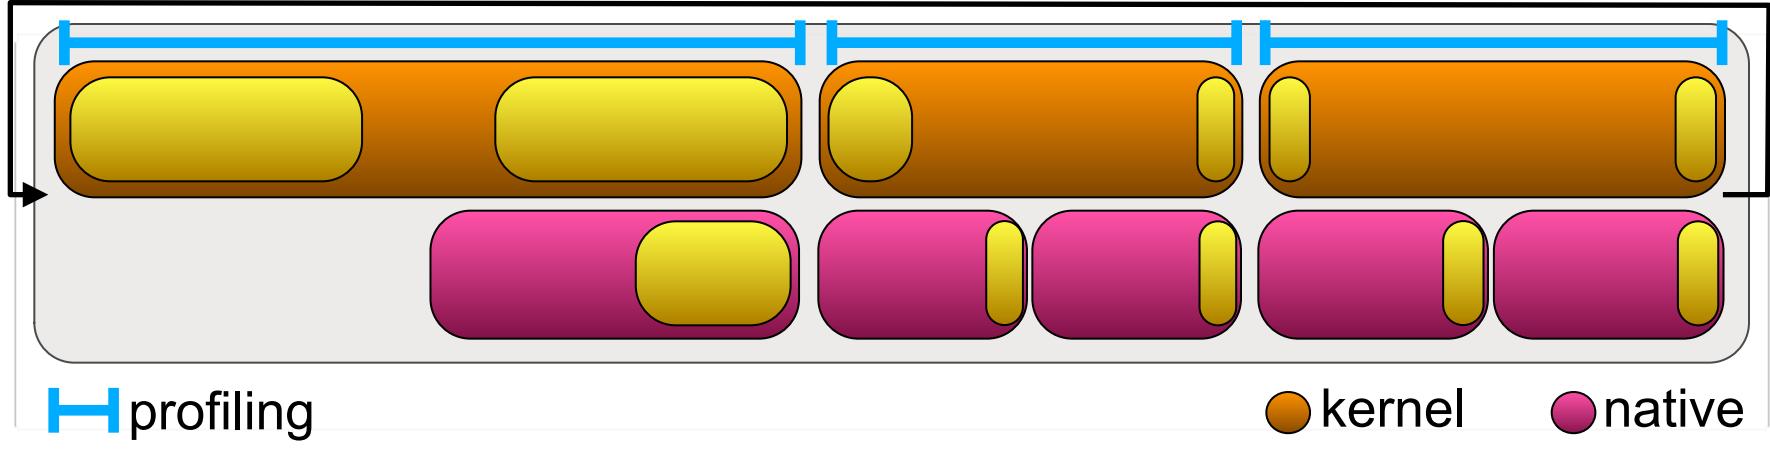


- **Allocate large fixed size memory buffers**
  - Use best estimate based on **kernel** execution time
  - Pass in data size as **kernel** argument
  - Adjust data size in **stream operation**



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Streaming Data



- **Dynamically adjust streaming data size**
  - Estimate reasonable time window for first **stream operation**
  - Use measured time window to adjust data size in **stream operation**



**SIGGRAPH ASIA 2009**  
the pulse of innovation

# OpenCL Event Model Usage

- **Summary**

- Use the execution model to your advantage
- Use synchronisation to your benefit
- Use events for fine grained control
- Plenty of flexibility already (even with immutable queues)



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# OpenCL Developer Insights

## • Know Your Platform

- Understand the cost of API calls
- Profile the overhead of your platform and commands
- Even compliant vendor implementations differ
- Give feedback to platform vendors
- Work with Khronos to improve OpenCL



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*

# Acknowledgements

- **Course Organisers**
  - AMD / NVIDIA / Khronos / Apple
- **University of Western Australia**
  - CMCA / WASP / iVEC

Derek Gerstmann

University of Western Australia

<http://local.wasp.uwa.edu.au/~derek>



**SIGGRAPH ASIA 2009**  
*the pulse of innovation*